Sampling

Data Analysis for Journalism and Political Communication (Fall 2024)

Prof. Bell

Definitions

- The group we are interested in studying is known as the **population**

Definitions

- The group we are interested in studying is known as the **population**
- Often, we are not able to count every unit in the population, so we take a **sample**

Definitions

- The group we are interested in studying is known as the **population**
- Often, we are not able to count every unit in the population, so we take a **sample**
- Our best guess about the population based on our sample is the **estimate**

Sampling

- There are many ways to derive an estimate from a sample, but recall that “garbage in = garbage out”: no amount of statistical wizardry can compensate for bad data

Sampling

- There are many ways to derive an estimate from a sample, but recall that “garbage in = garbage out”: no amount of statistical wizardry can compensate for bad data
- The key to any data analysis project is a quality sample, which is determined by two elements:

Sampling

- There are many ways to derive an estimate from a sample, but recall that “garbage in = garbage out”: no amount of statistical wizardry can compensate for bad data
- The key to any data analysis project is a quality sample, which is determined by two elements:
  1. A **random sample** of the population

Sampling

- There are many ways to derive an estimate from a sample, but recall that "garbage in = garbage out": no amount of statistical wizardry can compensate for bad data
- The key to any data analysis project is a quality sample, which is determined by two elements:
  1. A **random sample** of the population

Definition

The probability of any given unit being drawn from the population is uniform (the same)

Sampling

- There are many ways to derive an estimate from a sample, but recall that “garbage in = garbage out”: no amount of statistical wizardry can compensate for bad data
- The key to any data analysis project is a quality sample, which is determined by two elements:
  1. A **random sample** of the population
  2. The **sample size** is sufficiently large

In-class exercise

Sample Size

- How many units should you sample from the population?

Sample Size

- How many units should you sample from the population?
  It depends on your desired level of certainty.

Sample Size

- How many units should you sample from the population? **It depends on your desired level of certainty.**
- The most common level of certainty is 95% (the inverse of a **p-value** of .05, meaning that there is a 5% chance we are committing Type I error)

Sample Size

- How many units should you sample from the population? **It depends on your desired level of certainty.**
- The most common level of certainty is 95% (the inverse of a **p-value** of .05, meaning that there is a 5% chance we are committing Type I error)
- In other words, there is a 5% chance that the true population value is outside of the **confidence interval**

Sample Size

- How many units should you sample from the population? **It depends on your desired level of certainty.**
- The most common level of certainty is 95% (the inverse of a **p-value** of .05, meaning that there is a 5% chance we are committing Type I error)
- In other words, there is a 5% chance that the true population value is outside of the **confidence interval**
- If we re-sampled the population 100 times, 95 of our estimates would fall within the confidence interval (let's see this in action!)

Margin of Error

- There is a mathematical formula to estimate the confidence interval for a continuous mean, but more frequently we estimate a proportion

Margin of Error

- There is a mathematical formula to estimate the confidence interval for a continuous mean, but more frequently we estimate a proportion
- The confidence interval for a proportion is also called the **margin of error (MOE)**

Margin of Error

- There is a mathematical formula to estimate the confidence interval for a continuous mean, but more frequently we estimate a proportion
- The confidence interval for a proportion is also called the **margin of error (MOE)**
- The 95% MOE is calculated as:

$$1.96 * \sqrt{p * (1 - p)/n}$$

where p is the proportion and n is the sample size

Sample Size and the Margin of Error

- E.g. Washington Post-University of Maryland poll, December 14-18, 2023

Sample Size and the Margin of Error

- E.g. Washington Post-University of Maryland poll, December 14-18, 2023

$$1.96 * \sqrt{p * (1 - p)/n}$$

Sample Size and the Margin of Error

- E.g. Washington Post-University of Maryland poll, December 14-18, 2023

$$1.96 * \sqrt{p * (1 - p)/n}$$

$$1.96 * \sqrt{.33 * (1 - .33)/1024} = .029$$

Sample Size and the Margin of Error

- E.g. Washington Post-University of Maryland poll, December 14-18, 2023

  $$1.96 * \sqrt{p * (1 - p)/n}$$

  $$1.96 * \sqrt{.33 * (1 - .33)/1024} = .029$$

- We report the estimate with the MOE, i.e., 33 +/- 2.9%.

Sample Size and the Margin of Error

- E.g. Washington Post-University of Maryland poll, December 14-18, 2023

  $$1.96 * \sqrt{p * (1 - p)/n}$$

  $$1.96 * \sqrt{.33 * (1 - .33)/1024} = .029$$

- We report the estimate with the MOE, i.e., 33 +/- 2.9%.
- This means that there is a 5% chance that the true population proportion is outside of about 30-36%.

Sample Size and the Margin of Error

- E.g. Washington Post-University of Maryland poll, December 14-18, 2023

$$1.96 * \sqrt{p * (1 - p)/n}$$

$$1.96 * \sqrt{.33 * (1 - .33)/1024} = .029$$

- We report the estimate with the MOE, i.e., 33 +/- 2.9%.
- This means that there is a 5% chance that the true population proportion is outside of about 30-36%. *(Note: this does not mean that every value in the MOE is equally likely!)*

Sample Size and the Margin of Error

- E.g. Washington Post-University of Maryland poll, December 14-18, 2023

$$1.96 * \sqrt{p * (1 - p)/n}$$

$$1.96 * \sqrt{.33 * (1 - .33)/1024} = .029$$

- We report the estimate with the MOE, i.e., 33 +/- 2.9%.
- This means that there is a 5% chance that the true population proportion is outside of about 30-36%. *(Note: this does not mean that every value in the MOE is equally likely!)*
- Typically, pollsters will use a proportion (p) of .5 to calculate an MOE for the entire poll, rather than individual questions

Sample Size and the Margin of Error

- Mathematically, the MOE gets smaller as the sample size increases

Sample Size and the Margin of Error

- Mathematically, the MOE gets smaller as the sample size increases
- Intuitively, the larger our sample size, the less likely we are to draw an “outlier” sample and the more likely we are to get a representative sample

Sample Size and the Margin of Error

- Mathematically, the MOE gets smaller as the sample size increases
- Intuitively, the larger our sample size, the less likely we are to draw an “outlier” sample and the more likely we are to get a representative sample
- But the marginal improvement in the MOE from adding units to the sample decreases as the sample size grows

Sample Size and the Margin of Error

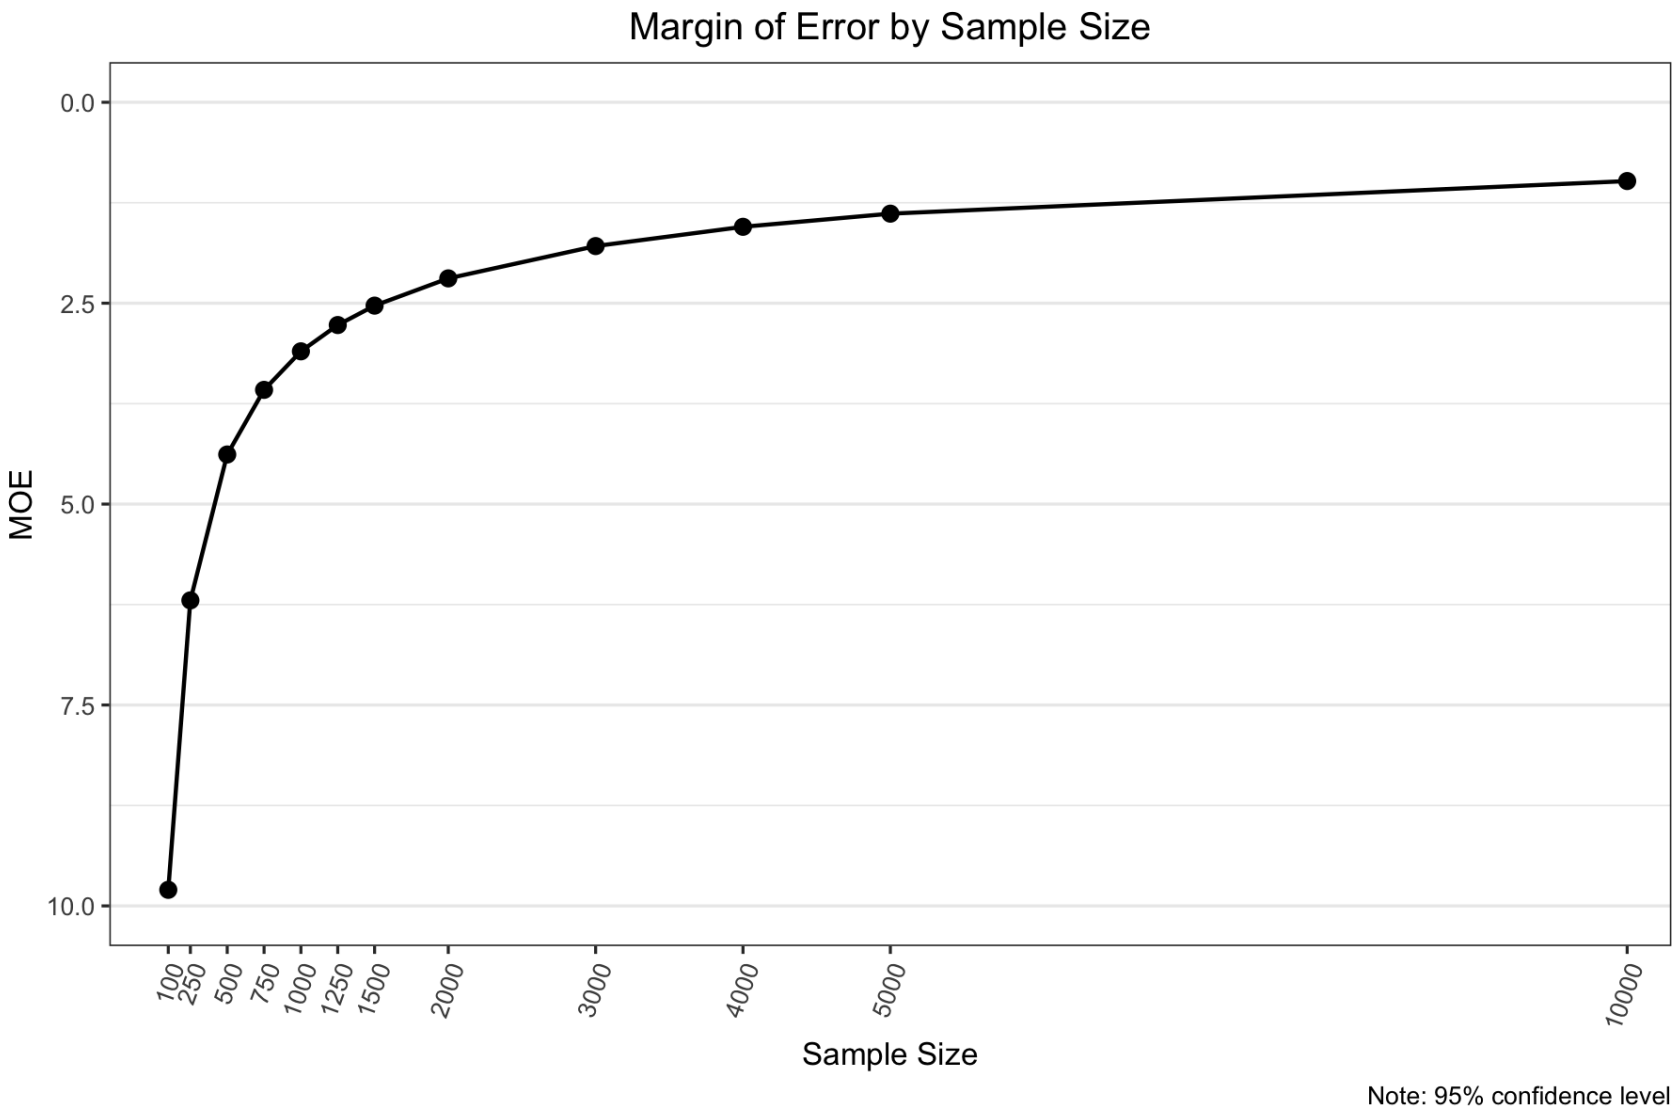

Sample Size and the Margin of Error

- Mathematically, the MOE gets smaller as the sample size increases
- Intuitively, the larger our sample size, the less likely we are to draw an "outlier" sample and the more likely we are to get a representative sample
- But the marginal improvement in the MOE from adding units to the sample decreases as the sample size grows
- Remember that the MOE only takes into account the sample size, not the potential for selection bias

Sample Size and the Margin of Error

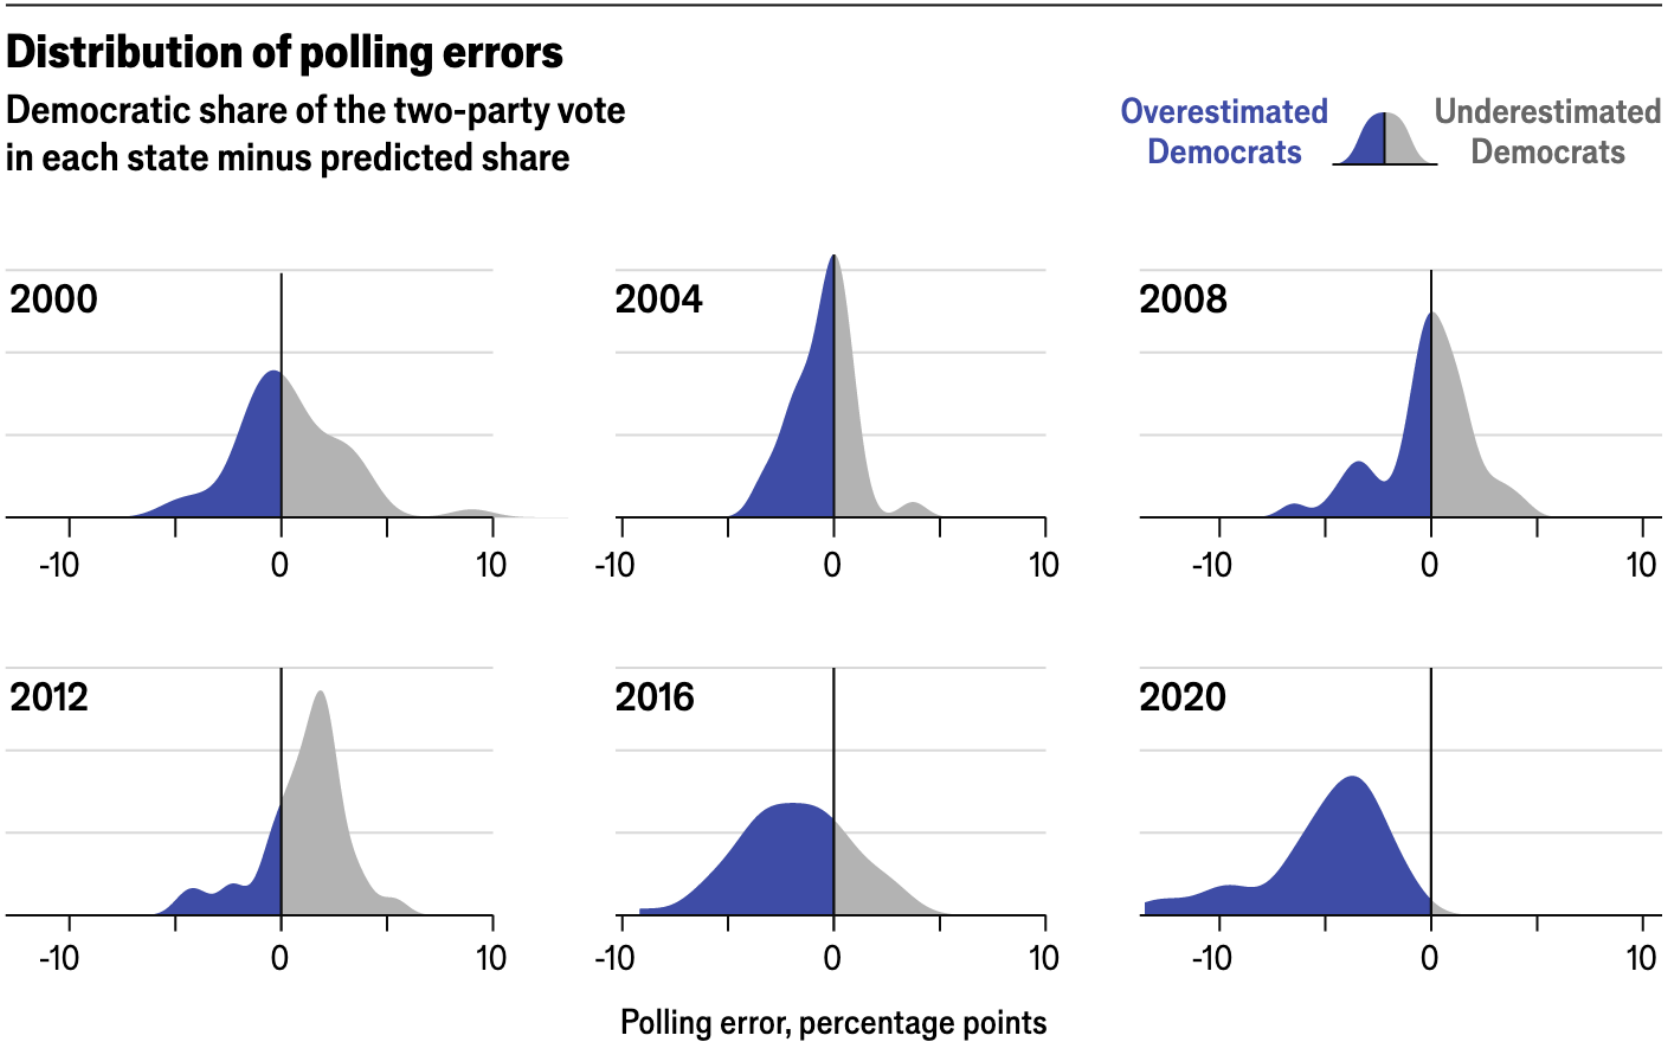

Source: The Economist

Getting a Random Sample

How would you generate a random sample of American voters?

Getting a Random Sample

- **Systematic Sampling**: Select every n^{th} unit of a population (e.g. exit polls)

Getting a Random Sample

- **Systematic Sampling**: Select every n^{th} unit of a population (e.g. exit polls)
- **List-based Sampling**: Randomly select units from an existing list of the population (e.g. registered voter lists)

Getting a Random Sample

- **Systematic Sampling**: Select every n^{th} unit of a population (e.g. exit polls)
- **List-based Sampling**: Randomly select units from an existing list of the population (e.g. registered voter lists)
- **Address-based Sampling (ABS)**: Randomly select households from a list of addresses provided by the U.S. Postal Service

Getting a Random Sample

- **Systematic Sampling**: Select every n^{th} unit of a population (e.g. exit polls)
- **List-based Sampling**: Randomly select units from an existing list of the population (e.g. registered voter lists)
- **Address-based Sampling (ABS)**: Randomly select households from a list of addresses provided by the U.S. Postal Service
- **Random-digit Dialing (RDD)**: Randomly select area codes, and then random digits are added to the end to create 10-digit phone numbers

Getting a Random Sample

- **Systematic Sampling**: Select every n^{th} unit of a population (e.g. exit polls)
- **List-based Sampling**: Randomly select units from an existing list of the population (e.g. registered voter lists)
- **Address-based Sampling (ABS)**: Randomly select households from a list of addresses provided by the U.S. Postal Service
- **Random-digit Dialing (RDD)**: Randomly select area codes, and then random digits are added to the end to create 10-digit phone numbers
- **Non-probability/Quota Sampling**: Pseudo-randomly selecting, from an opt-in pool of respondents, a sample that approximates the make-up of the general population